

Solution Cache: Remote Stores

Elton Stoneman
elton@sixeyed.com

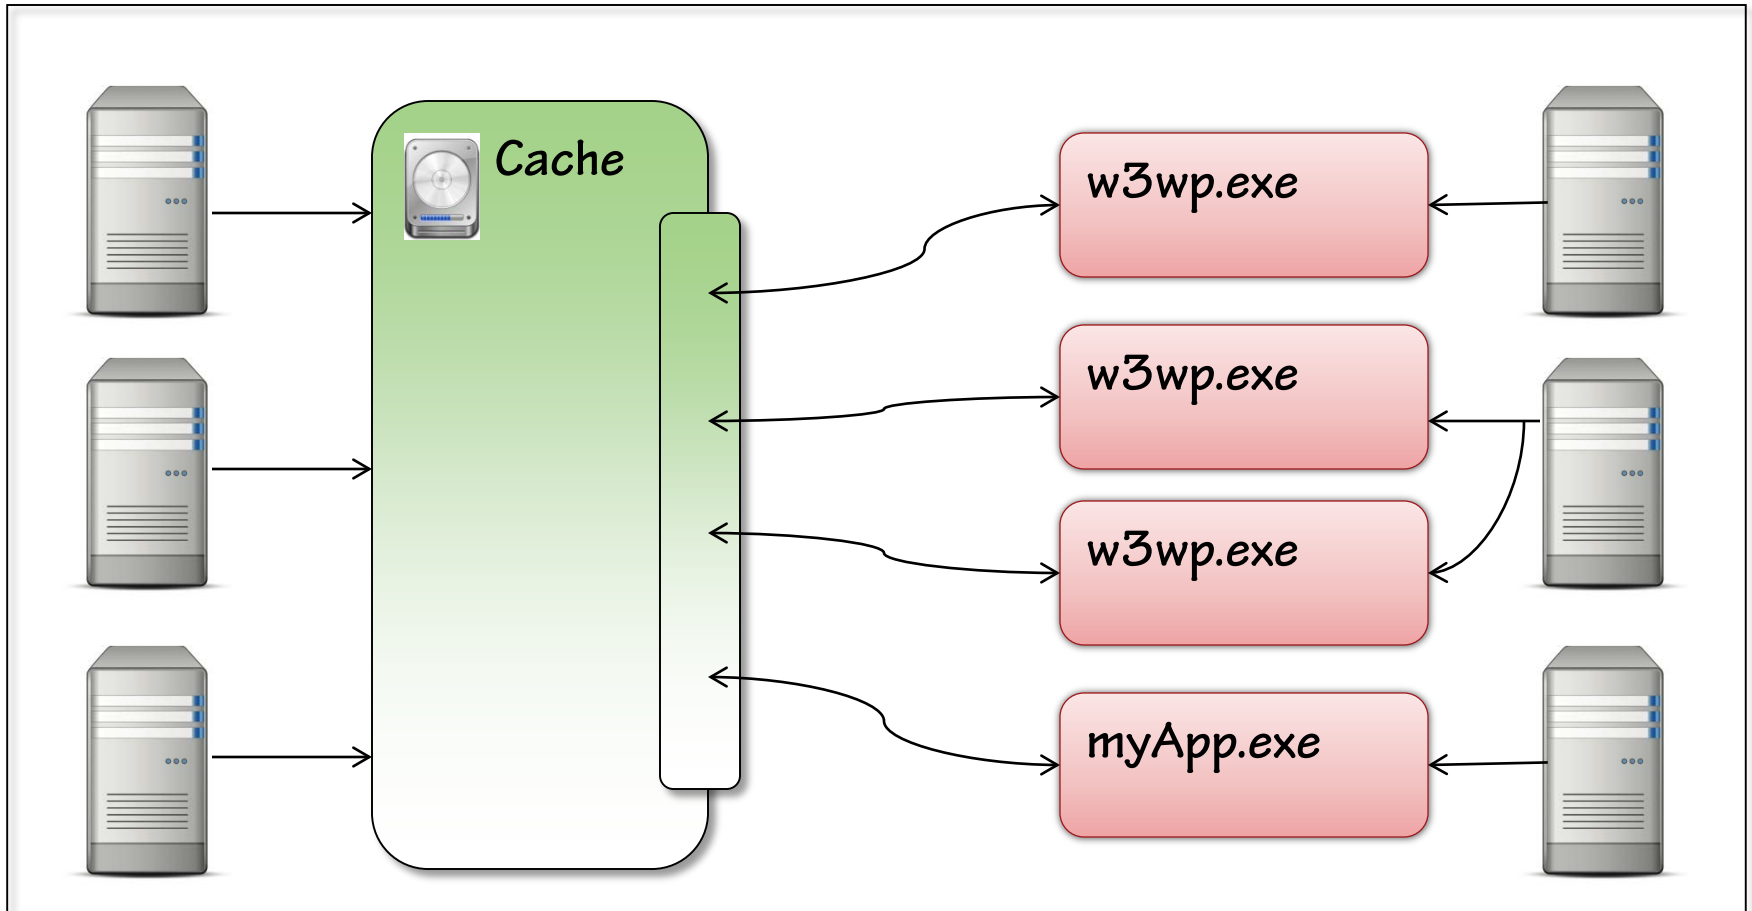


Outline

- **Remote cache stores**
 - Centralized
 - Distributed
 - Replicated
- **Remote caches**
 - Dependencies, usage and configuration
 - Management and extras
 - Matching against the decision matrix
- **Memcached**
- **Azure Table Storage**
- **DiskCache**
- **Applying the Decision Matrix**
 - Should I have multiple cache stores?

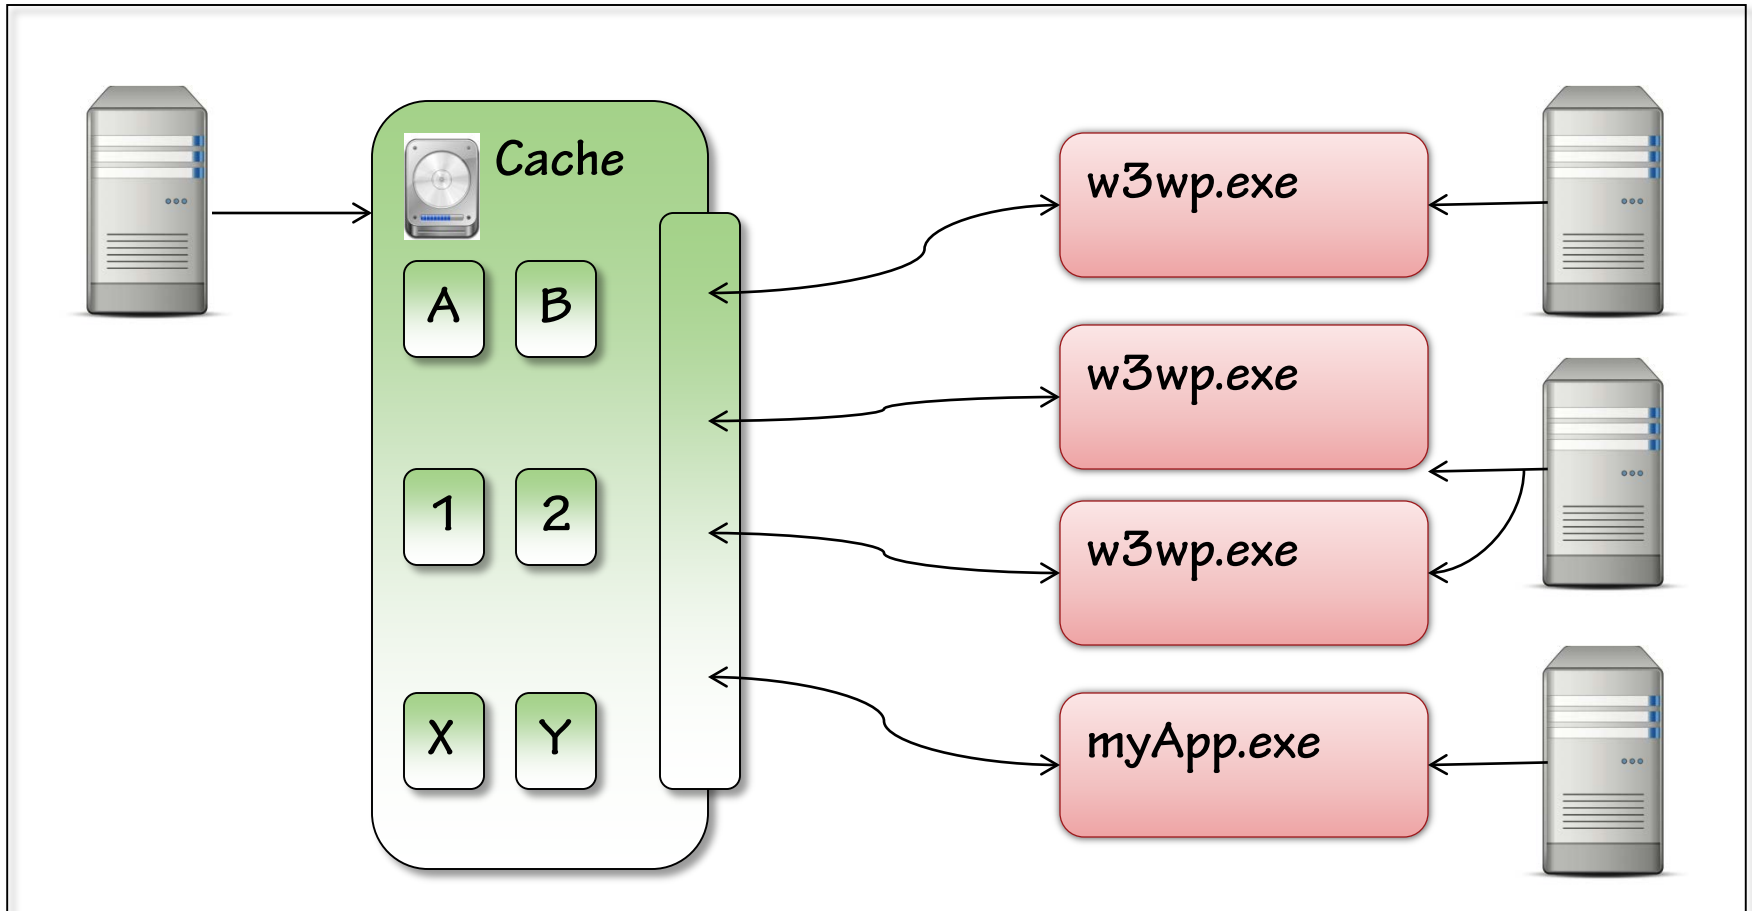
Remote Cache Stores

- Overview



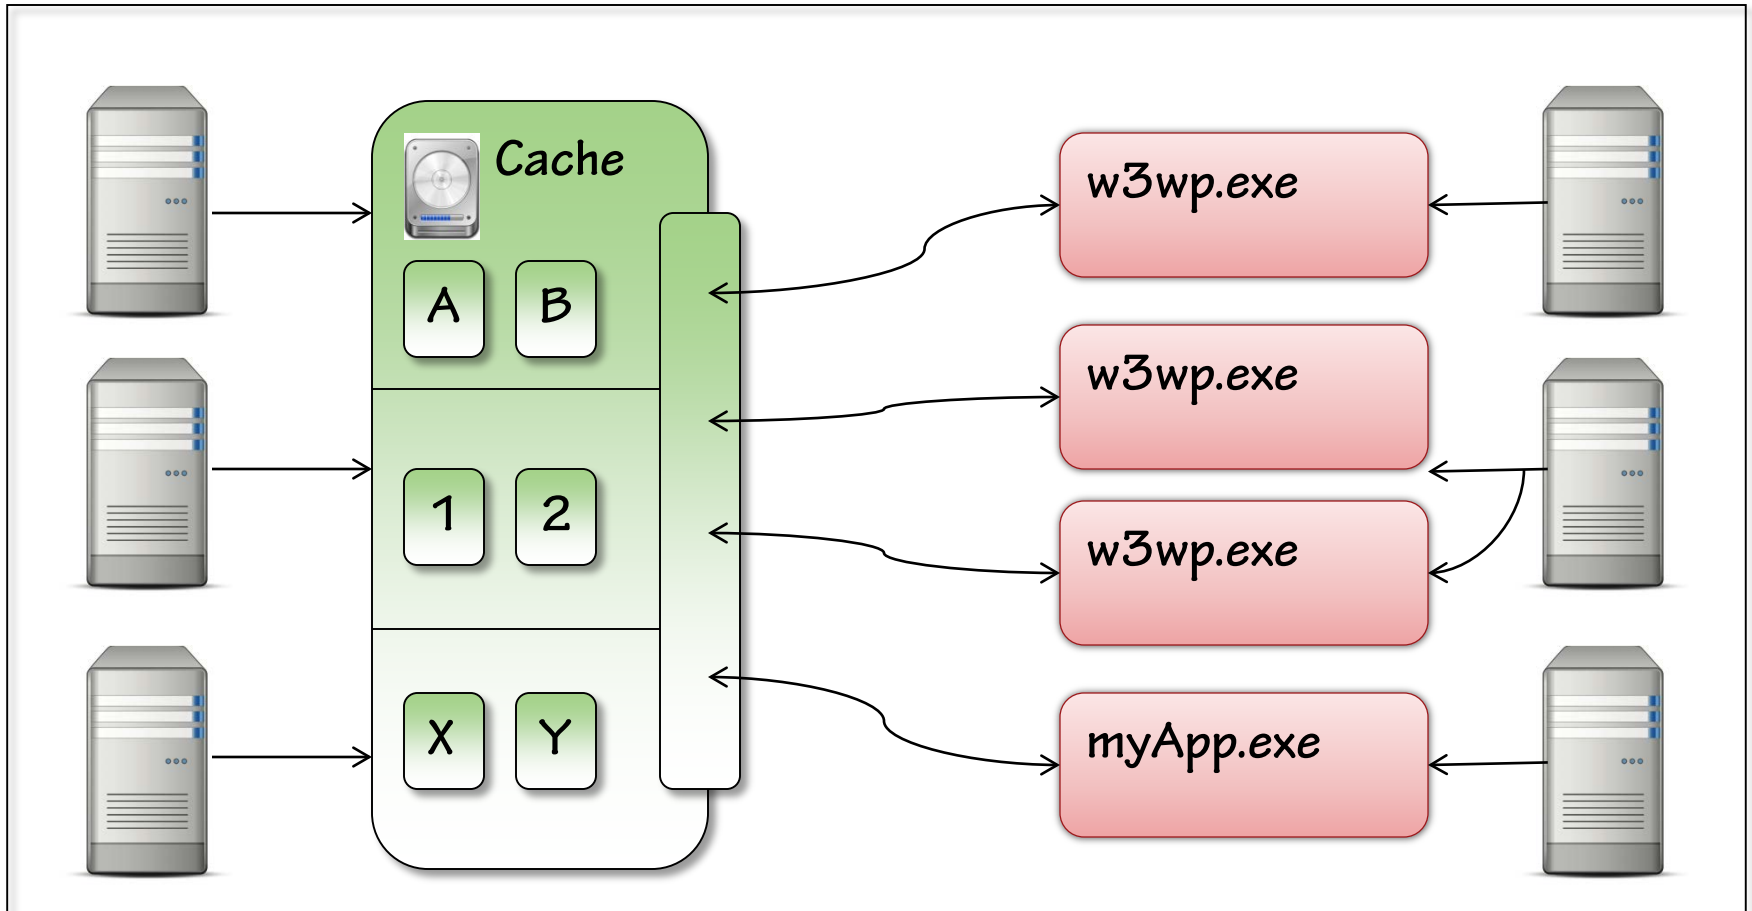
Remote Cache Stores

- Centralized



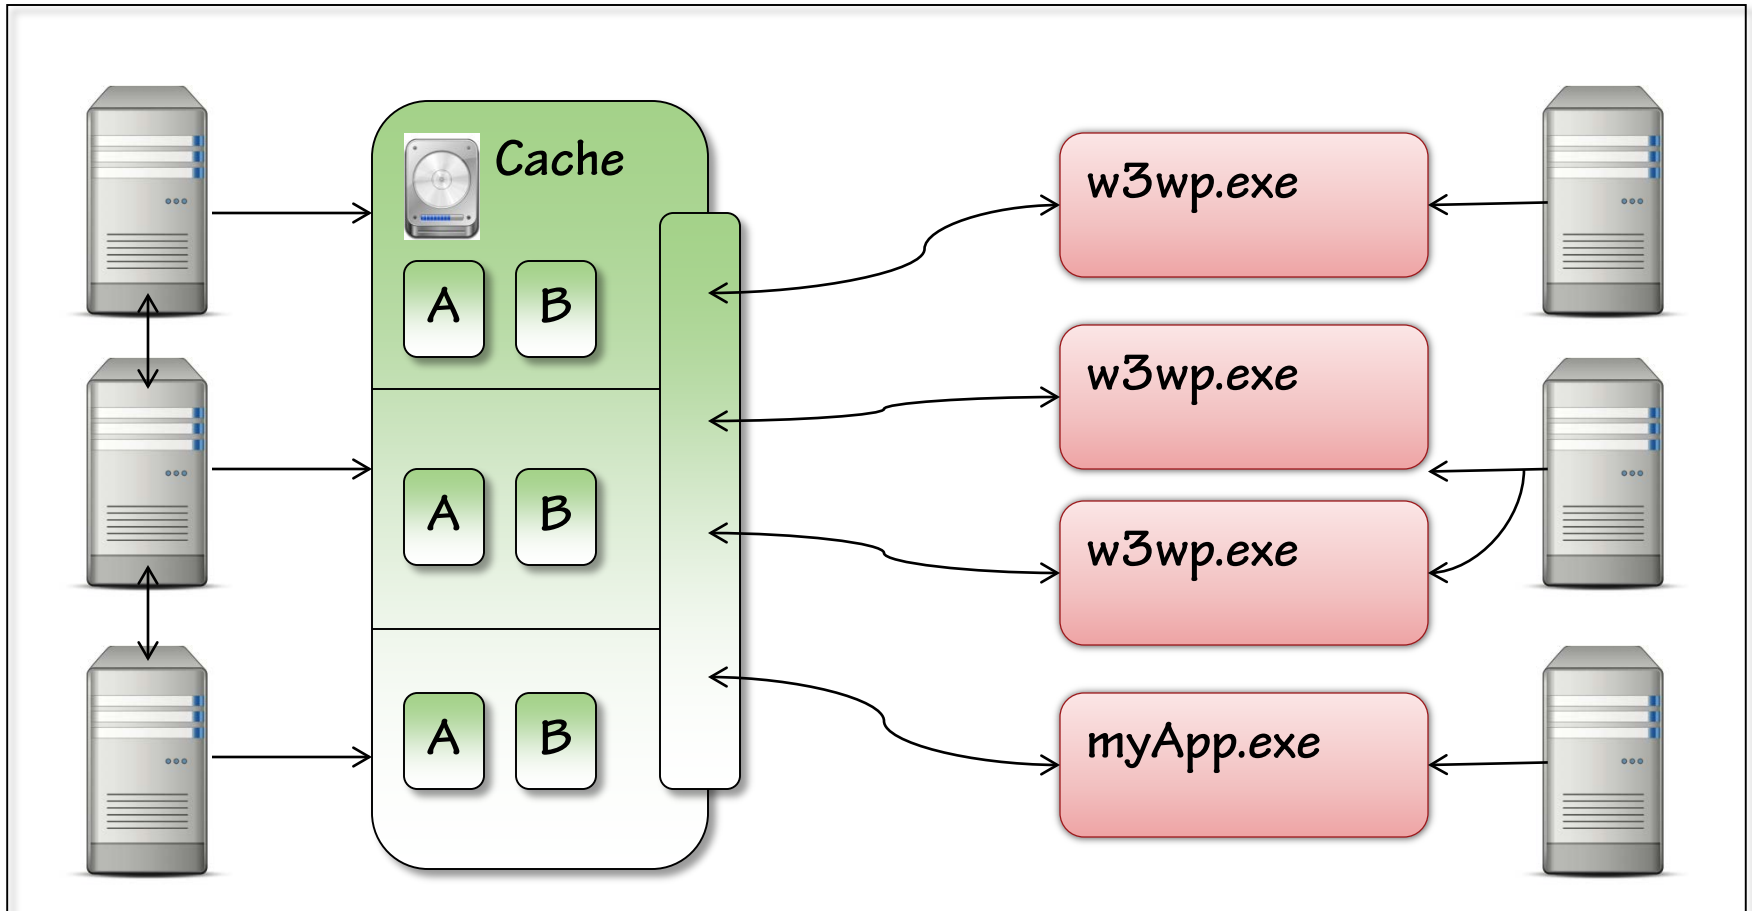
Remote Cache Stores

- Distributed



Remote Cache Stores

- Replicated



Memcached

- **Distributed cache**
- **Open source, cross-platform**
 - Evolved on Linux
 - Migrated to Windows with .NET client
 - Consistent protocol
- **Widely used**
 - Flickr, Twitter, Wikipedia, YouTube...

Memcached

- Demo

Memcached

- **Distributed cache**

- Runs in console or Windows Service
- Cache items split across nodes

- **Client configuration**

```
<enym.com>  
  <memcached protocol="Binary">  
    <servers>  
      <add address="127.0.0.1" port="11211" />  
    </servers>  
  </memcached>
```

- **No out-of-the-box management interface**

- Open protocol, 3rd party tools: Memcached Manager

- **Advanced features**


- Socket pool; performance monitors
- Protocol – AWS ElastiCache; EngineYard; Google App Engine

Memcached

- **Suitability applied to the decision matrix**
 - Cost – moderate, network access + memory lookup
 - Breadth – available to any process, any machine with network access
 - Longevity – machine uptime & network availability
 - Size
 - Total: sum of available memory on all nodes
 - Item: limited to 1Mb

Memcached

- Widely suitable, unless performance-critical



Azure Table Storage

- **Centralized cache**
 - PaaS "no-SQL" db used as store
- **Custom**
 - Simple implementation
 - Supports ICache and expiration
- **Tried and tested**
 - Extensible

Azure Table Storage

- Demo

Azure Table Storage

- **Centralized cache**
 - PaaS storage, shared logical node
 - NuGet client library
- **Connection string configuration**

```
<connectionStrings>  
  <!-- Azure Table Storage -->  
  <add name="Sixeyed.Core.Cache"  
        connectionString="DefaultEndpointsProtocol=https;  
                           AccountName=sccache;  
                           AccountKey=RW/LKHfeOBaYvjzIHZifXTsx..." data-bbox="147 436 722 595"/>
```


- **Management tools**
 - Azure Portal
 - Azure Storage Explorer
- **Extra features**
 - Persistence; geo-replication

Azure Table Storage

- **Suitability applied to the decision matrix**
 - Cost – slow, Internet access + disk IO
 - Breadth – available to any process, any machine with Internet access
 - Longevity – replication & network availability
 - Size
 - Total: Internet scale (100Tb)
 - Item: limited to 64Kb (1Mb with modification)

Azure Table Storage

- Good for smallish, long-lived items, balanced against performance



DiskCache

- **Centralized cache**
 - Items stored on network share
- **Custom**
 - Simple implementation
 - Supports ICache, cache size and expiration
- **Tried and tested**

DiskCache

- Demo

DiskCache

- **Centralized cache**
 - File system store
 - Single logical node
- **Simple configuration**

```
<sixeyed.core.caching>  
  <diskCache path="//192.168.2.160/cache"  
             encryptItems="false"  
             maxSizeInMb="200"/>  
</sixeyed.core.caching>
```


- **Familiar management tools**
 - Windows Explorer!
- **Extra features**
 - Encryption

DiskCache

- **Suitability applied to the decision matrix**
 - Cost – slow, network access + disk IO
 - Breadth – available to any process, any machine with network access
 - Longevity – disk lifespan & network availability
 - Size – available space on disk/SAN


DiskCache

- Good for long-lived items, balanced against performance



Applying the Decision Matrix

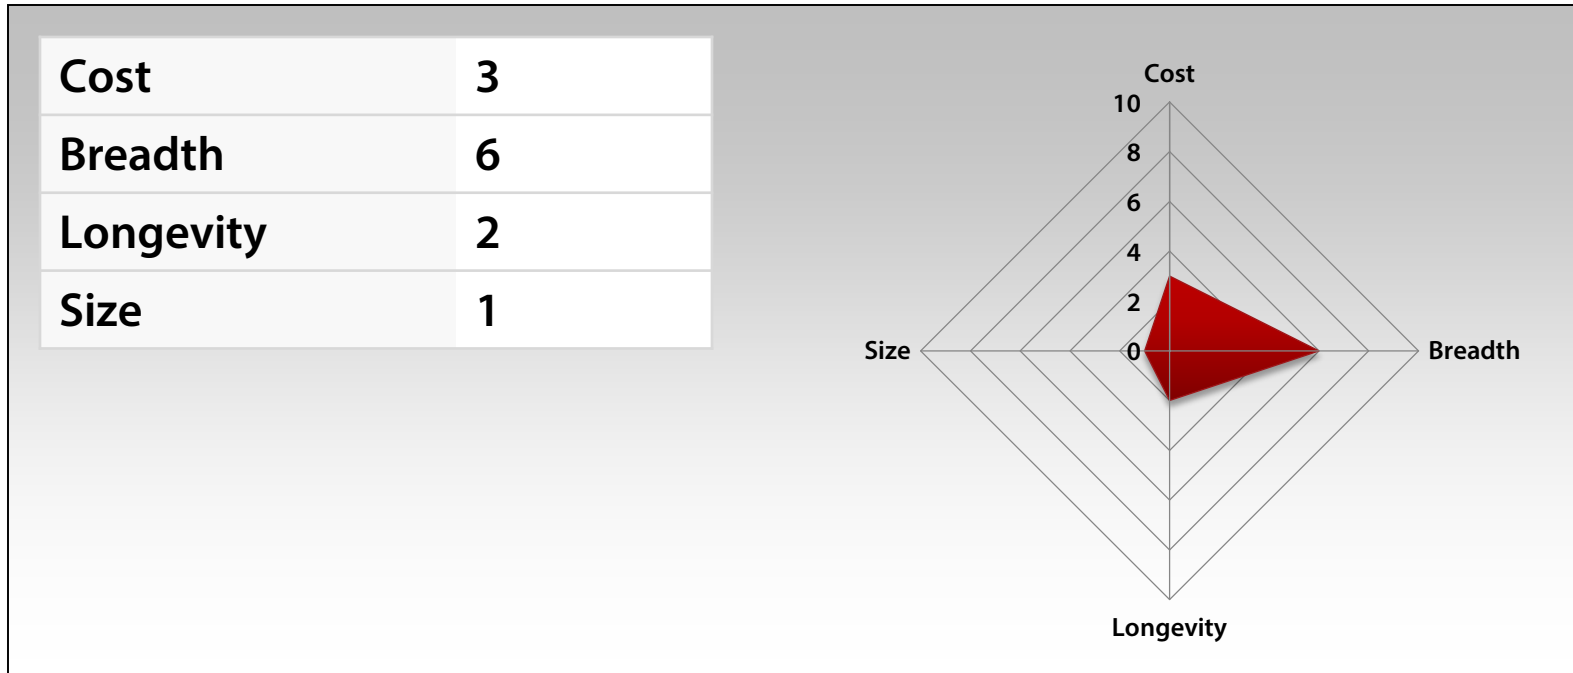
- Vehicle Makes



- Fast, accessible cache, reasonable longevity; size less important
 - Memcached
 - AppFabric Caching – replicated configuration

Applying the Decision Matrix


- CMS Content



- Fast, accessible cache; size and longevity less important
 - MemoryCache
 - NCache Express

Applying the Decision Matrix

- Quote Prices



- **Big, accessible, long-lived cache; speed less important**
 - DiskCache
 - Azure Table Storage

Applying the Decision Matrix

- Should I have multiple cache stores?

- Yes

- Three levels

1. Fast local memory store
2. Fast remote memory store
3. (Optional) Large persistent store

- Which ones?

	<i>Option 1</i>	<i>Option 2</i>
Level 1: fast local	.NET MemoryCache	NCache Express
Level 2: fast remote	AppFabric Caching	Memcached
Level 3: large persistent	DiskCache	Azure Table Storage

Summary

- **Types of remote cache store**
 - Centralized
 - Distributed
 - Replicated
- **Remote caches**
 - Memcached
 - Azure Table Storage
 - DiskCache
- **Applying the Decision Matrix**
 - Selecting the right cache
 - Balance performance with cost
 - Choosing your cache stores

References

■ Memcached

- Windows server – 32bit & 64bit
 - <http://s3.amazonaws.com/downloads.northscale.com/memcached-win32-1.4.4-14.zip>
 - <http://s3.amazonaws.com/downloads.northscale.com/memcached-win64-1.4.4-14.zip>
- .NET Client
 - <https://github.com/enyim/EnyimMemcached>
- Documentation
 - <https://code.google.com/p/memcached/wiki/NewStart>

■ Azure Table Storage

- Documentation
 - <http://msdn.microsoft.com/en-us/windowsserver/ee695849.aspx>